





Coalitions, Voting Power, and Computational Social Choice



Forming Coalitions: Coalitional Games ⁹

- *Coalitional games* model scenarios where agents *can benefit by cooperating*
 - Issues in coalitional games (Sandholm et al, 1999):
 - Coalition Structure Generation
 - Teamwork
 - Dividing the benefits of cooperation
-

Coalition Structure Generation

- Deciding *in principle* who will work together
- The basic question:
 - Which coalition should I join? 
- The result: *partitions* agents into disjoint *coalitions*
- The overall partition is a *coalition structure* 

Solving the optimization problem of each coalition

- Deciding how to work together.
 - Solving the “joint problem” of a coalition
 - Finding how to maximize the utility of the coalition itself.
 - Typically involves joint planning etc.
-

Dividing the Benefits

- Deciding “who gets what” in the payoff.
 - Coalition members cannot ignore each other’s preferences, because members can *defect*: if you try to give me a bad payoff, I can always walk away
 - We might want to consider issues such as *fairness* of the distribution
-

Formalizing Cooperative Scenarios

- A coalitional game:

$$\langle Ag, v \rangle$$

where

- $Ag = \{1, \dots, n\}$ is a set of agents
- $v = 2^{Ag} \rightarrow \mathbb{R}$ is the characteristic function of the game
- Usual interpretation: if $v(C) = k$, then coalition C can cooperate in such a way they will obtain utility k , which may then be distributed among the team members

Which Coalition Should I Join?

- Most important question in coalitional games:
is a coalition stable?
that is,
is it rational for all members of coalition to stay with the coalition, or could they benefit by defecting from it?
- (There is no point in me trying to join a coalition with you unless you want to form one with me, and vice versa.)
- Stability is a *necessary* but not *sufficient* condition for coalitions to form

The Core ⁹

- The *core* of a coalitional game is the set of *feasible* distributions of payoff to members of a coalition that *no* sub-coalition can reasonably object to
- An *outcome* for a coalition C in game $\langle \langle Ag, v \rangle \rangle$ is a vector of payoffs to members of C , $\langle \langle x_1, \dots, x_k \rangle \rangle$ which represents a *feasible distribution of payoff to members of Ag* .

Feasible means:

$$v(C) \geq \sum_{i \in C} x_i$$

Example

- If $v(\{1, 2\}) = 20$, then possible outcomes are $\langle 20, 0 \rangle$, $\langle 19, 1 \rangle$, $\langle 18, 2 \rangle$, \dots , $\langle 0, 20 \rangle$. (Actually there will be infinitely many!)

Objections

- Intuitively, a coalition C *objects* to an outcome if there is some outcome *for them* that makes *all of them* strictly better off
- Formally, $C \subseteq Ag$ objects to an outcome $\langle \langle x_1, \dots, x_n \rangle \rangle$ for the grand coalition if there is some outcome $\langle \langle x'_1, \dots, x'_k \rangle \rangle$ for C such that
$$x'_i > x_i \text{ for all } i \in C$$
- The idea is that an outcome is not going to happen if somebody objects to it!

The Core

- The *core* is the set of outcomes for the *grand coalition* to which *no* coalition objects
- If the core is *non-empty* then *the grand coalition is stable*, since nobody can benefit from defection
- Thus, asking ⁰
is the grand coalition stable?
is the same as asking:
is the core non-empty?

Problems with the Core

- Sometimes, the core is empty; what happens then?
- Sometimes it is non-empty but isn't "fair".

Suppose:

$Ag = \{1, 2\}$, $v(\{1\}) = 5$, $v(\{2\}) = 5$, $v(\{1, 2\}) = 20$.

Then outcome $\langle \langle 20, 0 \rangle \rangle$ (i.e., agent 1 gets everything) is *not* in the core, since the coalition $\{2\}$ can object. (He can work on his own and do better.)

However, outcome $\langle \langle 15, 5 \rangle \rangle$ *is* in the core: even though this seems unfair to agent 2, this agent has no objection.

- Why unfair? Because the agents are *identical*!

How To Share Benefits of Cooperation?

- The *Shapley value* is best known attempt to define how to divide benefits of cooperation fairly. It does this by taking into account *how much an agent contributes*.
- The Shapley value of agent i is the average amount that i is expected to contribute to a coalition.
- Axiomatically: a value which satisfies axioms: *symmetry*, *dummy player*, and *additivity*.

Shapley Defined

- Let $\delta_i(S)$ be the amount that i adds by joining $S \subseteq Ag$:
$$\delta_i(S) = \nu(S \cup \{i\}) - \nu(S)$$

...the *marginal contribution of i to S*

- Then the Shapley value for i , denoted

$$\text{Sh}(S, i) / \varphi_i = \frac{\sum_{r \in R} \delta_i(S_i(r))}{|Ag|!}$$

where R is the set of all orderings of Ag and $S_i(r)$ is the set of agents preceding i in ordering r .

Example: Shapley value

S	$v(S)$
$()$	0
(1)	1
(2)	3
(12)	6

Example: Shapley value

$$\begin{aligned} Sh(\{1, 2\}, 1) &= \frac{1}{2} \cdot (v(1) - v() + v(21) - v(2)) \\ &= \frac{1}{2} \cdot (1 - 0 + 6 - 3) = 2 \\ Sh(\{1, 2\}, 2) &= \frac{1}{2} \cdot (v(12) - v(1) + v(2) - v()) \\ &= \frac{1}{2} \cdot (6 - 1 + 3 - 0) = 4 \end{aligned}$$

Representing Coalitional Games

- It is important for an agent to know (for example) whether the core of a coalition is non-empty...so, how hard is it to decide this?
- Problem: naive, obvious representation of coalitional game is *exponential* in the size of Ag
- Now such a representation is:
 - utterly infeasible in practice; and
 - so large that it renders comparisons to this input size meaningless: stating that we have an algorithm that runs in (say) time *linear* in the size of such a representation means it runs in time *exponential* in the size of Ag

How to Represent Characteristic Functions?

- Two approaches to this problem:
 - try to find a *complete* representation that is succinct in “most” cases
 - try to find a representation that is not complete but is always succinct
 - A common approach:
interpret characteristic function over combinatorial structure


Representation 1: Induced Subgraph⁹

- Represent v as an undirected graph on Ag , with integer weights $w_{i,j}$ between nodes $i, j \in Ag$
- Value of coalition C then:


$$\nu(C) = \sum_{\{i,j\} \subseteq Ag} w_{i,j}$$

i.e., the value of a coalition $C \subseteq Ag$ is the weight of the subgraph induced by C

Representation 1: Induced Subgraph



the original graph defining v



subgraph induced by $\{A, B, C\}$
giving $v(\{A, B, C\}) = 3 + 2 = 5$

Representation 2: Marginal Contribution Nets ⁹

(leong & Shoham, 2005)

- Characteristic function represented as rules:
pattern \rightarrow value

Representation 2: Marginal Contribution Nets

- Pattern is conjunction of agents, a rule *applies* to a group of agents C if C is a superset of the agents in the pattern
- Value of a coalition is then sum over the values of all the rules that apply to the coalition
- Example:
$$\begin{array}{l} a \wedge b \longrightarrow 5 \\ b \longrightarrow 2 \end{array}$$

We have: $v(\{a\}) = 0$, $v(\{b\}) = 2$, and $v(\{a, b\}) = 7$

- We can also allow negations in rules (agent not present)

Example: Marginal Contribution Nets

Consider the marginal contribution net:

a	\longrightarrow	2
$a \wedge b$	\longrightarrow	7
b	\longrightarrow	3
c	\longrightarrow	4
$b \wedge c$	\longrightarrow	-3

(a) Let ν be the characteristic function defined by these rules. Give the values of the following:

1. $\nu(\emptyset)$
2. $\nu(\{a\})$
3. $\nu(\{b\})$
4. $\nu(\{a, b\})$
5. $\nu(\{a, b, c\})$

Example: Marginal Contribution Nets

Answer:

$$\nu(\emptyset) = 0, \nu(\{a\}) = 2, \nu(\{b\}) = 3, \nu(\{a, b\}) = 12, \nu(\{a, b, c\}) = 13.$$